



Audit and Upgrade Plan for the GENESIS Trading Bot

Overview and Current Architecture

GENESIS is a MetaTrader 5 trading bot with a Python backend and a TypeScript/React frontend. It currently follows an *event-driven microservices* architecture where each subsystem (data feed, signal processing, execution, risk management, etc.) runs independently and communicates through a high-throughput event bus (e.g. Kafka or Redis Streams) ¹. This design provides institutional-grade performance – all services consume a time-ordered stream of events (ticks, signals, news, commands) and maintain synchronized state ². The decoupled, asynchronous communication yields **low latency and fault tolerance**: a module crash won't bring down the whole system, and events can be replayed from the log for recovery or backtesting ³. Each service is containerized for independent deployment and scaling ⁴. The overall goal for GENESIS is to be a **"sniper-style" trading system** – adaptive, resilient, and fully audit-ready (every decision and action is logged) while always protecting capital ⁵. Crucially, the bot is designed to be **FTMO-compliant** from the ground up, respecting strict rules (max 5% daily loss, 10% total drawdown, 1:30 leverage, etc.) as hard constraints in its logic ⁶.

Despite this solid foundation, our audit has identified **key limitations and improvement areas** in performance, risk controls, and feature integration. Below, we present a comprehensive upgrade plan addressing these issues, along with recommendations for architecture enhancements, stricter FTMO-aligned risk management, news/alert integrations, containerized deployment for multi-account use, and overall code quality improvements. We then outline a prioritized action plan with expected outcomes for each step.

1. Full-Stack Upgrade Plan: Performance, Latency, Fault-Tolerance & Execution

To achieve a 10x improvement in performance and reliability, we propose several upgrades across the stack:

- **High-Performance Event Bus & Async Processing:** Ensure the event-driven design is fully leveraged for throughput and low latency. The Market Data Feed service should use asynchronous I/O or dedicated threads to read ticks in a tight loop, so each price tick is published to the bus within milliseconds ⁷ ⁸. This will minimize latency from market to bot. We recommend using Kafka (or Redis Streams) as the central message broker in production for its durability and scale – Kafka can persist events so services can restart and catch up from the last event, and a Kafka cluster removes single points of failure ⁹. The system already embraces asynchronous, non-blocking communication to avoid bottlenecks ¹⁰; we will audit all modules to eliminate any remaining synchronous or CPU-bound chokepoints. This may include optimizing indicator calculations (using vectorized libraries or C++ extensions if needed) and introducing concurrency (e.g. thread pools or `asyncio` tasks) for tasks that can run in parallel. The expected outcome is **lower end-to-end**

latency (from tick to trade execution) and higher throughput, allowing GENESIS to handle more instruments and higher tick rates without lag. We will continuously monitor latency (e.g. timestamp each tick vs when it's processed) to verify improvements ¹¹ .

- **Optimized Data Feed with Redundancy:** Upgrade the data ingestion layer for robustness and speed. GENESIS should maintain **redundant price feeds** – e.g. primary connection to MT5 and a secondary feed (another broker API or a FIX data feed) for failover ¹² . If the primary feed stalls or disconnects, the system can seamlessly switch to the backup to avoid data gaps ¹³ . We will also implement sequence checking to ensure ticks are processed in correct order and detect any gaps/out-of-order data ¹⁴ . If any missing ticks or feed pauses occur, the system will log alerts and attempt reconnection immediately ¹³ . By buffering incoming ticks and recovering on reconnect, we prevent data loss during feed hiccups. This redundancy and fast reconnect logic greatly improve **fault-tolerance** in live trading – a temporary broker API outage won't blind the strategy. Additionally, the feed module should run health checks (e.g. “no tick for > X seconds” triggers an alert) to promptly catch feed issues ¹⁵ . These measures ensure continuous, reliable data flow to all other services.
- **Microservice Hardening & Fault Isolation:** We will reinforce the microservice boundaries to improve resilience. Each core component (Market Data, Signal Engine, Execution Engine, Risk Engine, etc.) will be run in its own process/container with lightweight *health checks* and auto-restart policies ¹⁶ . If any service encounters an error, it can be restarted independently (via Docker or an orchestrator) without bringing down the whole bot ³ . We'll implement comprehensive exception handling around each service's main loop – e.g. wrapping processing loops in try/except to catch any unexpected exceptions, log them with stack traces, and keep the service running safely ¹⁷ . In case of a fatal error, the system will enter a **safe fail state**: for example, if the ExecutionEngine experiences a serious error, it should set a global flag (or notify the RiskEngine) to halt new trading until the issue is resolved ¹⁸ . Similarly, if one service crashes entirely, the orchestrator (Docker Compose or Kubernetes) will automatically restart it, and a supervisor or the Dashboard will raise an alert to operators ¹⁹ . These fault-tolerance enhancements guarantee **24/7 uptime**: the bot can recover from component failures gracefully, and no single bug will cascade into a full system shutdown.
- **“Sniper-Style” Trade Execution Improvements:** A major upgrade focus is the execution logic, to achieve precise, low-slippage entries and exits. Going forward, **market orders will be avoided entirely** – the Execution Engine should use *limit orders only* for all entries and exits ²⁰ . This sniper-style tactic ensures we control the fill price and eliminate negative slippage at entry. Concretely, for a buy signal, GENESIS will place a limit order a few pips *below* the current ask (for a sell, a few pips above the bid) so that we only get filled at a better price than the current market ²¹ . This way, if the market moves favorably into our limit, we get a price improvement; if not, we simply don't trade rather than chase a poor price. By guaranteeing a price or better, limit orders enforce discipline (at the risk of the order not filling) ²¹ . To further refine execution, we'll implement **adaptive order placement** logic: if a limit order isn't filled within a short window, the bot can *re-quote* once or twice – i.e. adjust the limit price closer to market – to attempt a fill, but with a strict cap on how far it will chase ²² ²³ . This prevents the bot from following a runaway price and entering too late. Any unfilled orders after the allowed retries will be aborted to avoid slippage. We'll also handle **partial fills** gracefully: for example, if we send an order for 1.0 lots and only 0.5 lots execute before price moves, the ExecutionEngine will detect the partial fill via MT5 order updates ²⁴ . It can then cancel

the remaining volume and decide whether to re-place it (perhaps at a new price) or forego it if conditions have changed. Often splitting the order into smaller chunks or waiting for a brief pullback can help fill the remainder ²⁵. We'll implement a short retry loop for partial fills (e.g. try to fill the leftover for a few seconds with new limit orders) and then give up if not filled ²⁵. These sniper execution upgrades will significantly **reduce slippage and improve fill quality**, essentially behaving like a Smart Order Router. *(Note: On MT5, we are typically limited to one liquidity provider, but if multiple accounts or an aggregator API were available, the system could even split orders across venues to get best price ²⁶.)* In addition, we'll enforce that **every order has an immediate stop-loss attached** – the ExecutionEngine will always transmit orders with a linked SL level, as MT5 allows ²⁷. This is not only an FTMO requirement (no trade without SL) but also protects every position from catastrophic loss the moment it's opened. Overall, by upgrading to this "surgical" execution approach (precise limit entries, quick adjustments, mandatory SL), we expect to see better entry prices on average, controlled risk on each trade, and adherence to FTMO's no-market-order preference.

- **Additional Execution Enhancements:** We will incorporate a few more best practices to bolster execution effectiveness. The ExecutionEngine will include **higher-timeframe confirmation checks** just before final order placement – for instance, double-check the 4H or 1D trend is still aligned with the trade direction at the moment of execution; if a new higher-timeframe candle just flipped against our signal, the engine may abort the entry as conditions have changed ²⁸. This prevents acting on signals that were valid moments ago but have since become stale or conflicted. We'll also implement **time-of-day filters** in the ExecutionEngine: the bot will avoid trading during known poor liquidity periods (for example, the low-liquidity rollover around 5pm New York time when spreads widen) ²⁹. We can configure a schedule (e.g. disable new trades between 4:30pm–6:00pm NY, or other quiet periods) to reduce avoidable slippage. Likewise, we can restrict trading to major market sessions (London/New York overlap, etc.) when liquidity is highest ²⁹. These execution-time checks and filters act as a "sniper's intuition," adding another layer of protection so trades are only executed under optimal conditions. The expected outcome is a further improvement in trade performance metrics: higher average return-to-risk per trade due to better entries, and fewer losses from catching bad trades in illiquid moments.
- **Scalability and Co-location:** As a longer-term upgrade, we recommend deploying GENESIS on high-performance infrastructure close to the broker's servers (for minimal network latency). For example, if FTMO's liquidity is in London, running the bot on a London-based VPS or cloud region will cut down order transit times. The microservice architecture already supports horizontal scaling – if needed, we can scale out certain services (like running multiple instances of the SignalEngine consuming the same feed in a consumer group) to handle load ⁴ ²⁹. The event-bus design inherently allows adding more consumers without altering producers. Thus, as the strategy complexity or number of instruments grows, GENESIS can scale by simply adding resources rather than refactoring core logic. This positions the bot for **institutional-grade scalability**. In summary, the full-stack upgrades above (faster async processing, robust failovers, and improved execution tactics) collectively aim to make GENESIS both **fast and resilient** – ensuring it captures opportunities in real-time while gracefully handling errors or volatility spikes.

2. Stricter FTMO-Compliant Risk Management and Execution Controls

FTMO's rules are unforgiving – a breach of the daily loss or total drawdown limits means instant failure of the challenge. Therefore, GENESIS must **rigidly enforce these risk constraints in real-time**, with automatic kill-switches and pre-trade checks to prevent any violation. We will bolster the existing RiskEngine and ExecutionEngine with the following FTMO-aligned controls:

- **Maximum Daily Loss (5%) Kill-Switch:** The bot will continuously track the day's running P/L (including both closed trades and current open unrealized losses) against the allowed 5% drawdown from the starting balance. If the day's loss reaches **around 90% of the limit** (e.g. -4.5% equity drawdown), GENESIS will trigger a **preemptive kill-switch**: no new trades will be allowed for the rest of that day ³⁰. This provides a safety buffer before hitting the 5% threshold. If, despite this, the loss *does* hit the full -5% limit intraday, the system will immediately **close all open positions** to stop the bleeding and halt all further trading for that day ³¹. Essentially, the trading day is considered over if -5% is reached ³². These rules turn FTMO's guideline into hard-coded behavior: the bot will not "accidentally" violate the Max Daily Loss because it will cease trading at ~4.5% drawdown and forcibly flatten the account at 5%. Each new trading day at **UTC midnight**, the daily loss counter resets to zero (since FTMO resets the limit each day) ³³. We will implement this reset logic on a schedule so the bot knows a new day's limit is fresh (taking care to align with FTMO's timezone cutoff) ³⁴. With this kill-switch in place, the **expected outcome is elimination of catastrophic down days** – the worst-case daily loss is capped just under 5%, and the bot lives to fight another day.
- **Maximum Total Drawdown (10%) Protection:** Similarly, the RiskEngine will monitor the **overall drawdown** from the initial account balance (e.g. for a \$200k account, -\$20k is the 10% max loss). If equity falls to that 10% drawdown level at any time, trading must halt. In practice, we will again stop out slightly earlier for safety – e.g. at 9% or so – but certainly by 10%, the system will **halt all trading entirely** ³⁵. Exceeding 10% is an instant failure of the FTMO challenge, so our bot should never allow reaching that point ³⁶. If the total drawdown limit is tripped, GENESIS will close any open trades and not execute further until manually reset (since an FTMO challenge would be over). These absolute loss limits ensure the bot **never violates the FTMO loss rules**; at worst, it will stop trading and preserve the account from further losses. (We will make the reference balance for drawdown calculations configurable, because in FTMO's Verification or funded accounts the limits might reset or use a new balance – the system can be told what baseline to use.)
- **Leverage Cap (1:30) Enforcement:** FTMO mandates maximum leverage of 1:30, meaning at any time margin usage cannot exceed roughly 3.33% of account equity per trade (since $1/30 \approx 0.0333$). The RiskEngine will perform a **pre-trade margin check** for every order: it will calculate the required margin for the position size (based on the symbol's contract size and current price) and ensure that adding this trade will not push the account beyond the 1:30 leverage ratio ³⁷. In practical terms, we will check that `position_size * contract_size / equity <= 30` (or equivalently margin used $\leq 3.33\%$ of equity) ³⁷. If a proposed trade would exceed the leverage cap, the trade is blocked immediately. This is done **before** sending the order to MT5. By simulating the worst-case margin impact of a new order, we guarantee the account stays under the leverage limit at all times. Typically, FTMO's 30x leverage on a \$200k account means we can utilize up to ~\$6.6k margin at once; our risk checks will strictly honor that. The expected outcome is that GENESIS will never breach the margin

requirements – any oversize order is simply rejected with a logged warning (and perhaps an alert to notify the trader that position sizing was too large). This prevents accidental over-leveraging, which is especially important if the strategy trades multiple instruments; the total exposure across all should be monitored. We'll also implement **position correlation limits**: for example, if one trade is open on EURUSD, the bot might restrict opening another highly correlated pair like GBPUSD with full size, to avoid effectively doubling exposure on the USD ³⁸. Such rules further ensure the bot stays within a conservative risk envelope compliant with FTMO's spirit.

- **No Market Orders – Limit-Only Execution:** As mentioned earlier, we are banning market orders entirely in favor of limit orders for entries and exits. This will now be an **enforced rule in the ExecutionEngine**. We will modify the code such that any attempt to send a market order either gets converted to a limit order at a sensible price or is outright refused ²⁰. By doing so, we uphold a “no market order” policy that aligns with a disciplined trading approach (FTMO doesn't explicitly forbid market orders, but this is a self-imposed rule to control slippage). The ExecutionEngine's Copilot prompt explicitly calls for only using limit orders to control price ²⁰, and we have implemented that sniper logic above. Thus, this item is essentially covered: **every trade will be a limit or stop-limit order** with predefined entry price and stop-loss. The benefit is that we never enter trades at arbitrary prices – it's always at a planned level or not at all – which protects us from slippage-induced drawdowns that could jeopardize the daily loss limit. Combined with the above risk checks, this ensures a very tight compliance: the bot knows the maximum loss of each trade upfront (from the stop-loss and size) and ensures even that worst-case won't break FTMO rules ³⁹ ⁴⁰.
- **Mandatory Stop-Loss on Every Trade:** FTMO requires that *every* position must have an initial stop-loss in place. We will enforce **no trade without a stop-loss** as an immutable rule. In practice, our ExecutionEngine already attaches a stop-loss to every order when sending it (MT5 allows including SL/TP in the order placement) ²⁷. We will add safeguards such that if for any reason a signal comes through without a suggested SL, the trade is rejected – “no SL = no trade” ⁴¹. The RiskEngine can double-check after the fact as well: if it ever detects an open position without a stop, it should immediately either apply a default stop or close the position and log an error. This double layer ensures we **never have an unprotected position**. The stop-loss levels will be chosen by the strategy (e.g. based on ATR or support/resistance), but enforced by the system.
- **Conservative Risk per Trade & Session:** In addition to hard FTMO limits, we will maintain conservative risk metrics internally. For example, position sizing per trade will target maybe 0.5%–1% risk (so that even a full SL hit won't drop the account more than 1%). We'll also consider adding **dynamic kill-switches** beyond FTMO's rules, as best practices. These include: a *consecutive losses halt* – if the bot hits, say, 3 losses in a row, it pauses trading for a “cool-off” period to avoid tilt trading ⁴²; a *volatility spike halt* – using an indicator like ATR or an external volatility index, if market volatility jumps to extreme levels (beyond what the strategy was calibrated for), pause new entries ⁴³; and a *news shock halt* – if an unscheduled major news event occurs (detected via the NewsProcessor, see next section), immediately freeze trading because such events can invalidate signals ⁴⁴. These layered protective switches can all flip a master `pause_trading` flag that the ExecutionEngine and SignalEngine honor ⁴⁵. While not strictly mandated by FTMO, they provide an extra safety net, containing risk in unusual conditions. The **expected outcome** of all the above risk controls is a bot that is extremely safe: it will practically **impossible for it to violate FTMO rules** or take outsized losses without human override. All potential breaches are pre-empted (blocked at the

order level or halted via kill-switch). This not only preserves the challenge account but also instills discipline akin to a professional risk manager overseeing the algorithm.

- **Detailed Risk Logging and Journaling:** Compliance is not only about rules but also about transparency. We will enhance the journaling of all risk-related decisions. For every trade signal, the RiskEngine will log whether it was allowed or blocked and why. For instance: *“Signal BUY 1.5 lots EURUSD blocked – would exceed daily loss limit (projected -5.2%).”* Such logs give an audit trail that our risk checks are functioning ⁴⁶. We will also maintain a running tally of daily P/L in the logs, so one can see how close to limits we are at any point. If a kill-switch triggers, it will immediately log an alert like *“Kill-switch activated: Daily loss threshold reached, trading halted.”* and also publish an event that can be picked up by the alerting system (to notify humans). These measures ensure that if FTMO or an auditor reviewed our bot, we could demonstrate every precaution we took to follow the rules. In sum, GENESIS will behave as if an FTMO risk officer is embedded in its code – never allowing forbidden actions and always keeping meticulous records.

3. Integrations for News, Alerts, Journaling, and Multi-Account Orchestration

Modern swing trading requires situational awareness beyond raw price data. We propose integrating external information feeds and alerting mechanisms to make GENESIS more intelligent and user-friendly, as well as preparing it to manage multiple accounts in parallel.

- **Real-Time News and Event Awareness:** Volatility around news events can make or break a trading day. GENESIS should proactively **monitor economic news and announcements** and adjust its trading accordingly. We will implement a **NewsProcessor** microservice that aggregates economic calendars and real-time news feeds. For scheduled events, we’ll pull data from a source like **ForexFactory’s API or an economic calendar feed** to know when high-impact events (e.g. central bank rate decisions, CPI releases) are coming ⁴⁷ ⁴⁸. The bot will automatically pause new trades starting e.g. 15–30 minutes before a *High impact* event and not resume until a certain time after, to avoid whipsaw volatility ⁴⁹. For real-time breaking news, we can use APIs from services like **Reuters or Bloomberg** (if available) or other financial news APIs. The NewsProcessor will continuously fetch headlines and news bulletins, filtering for market-moving keywords or symbols of interest ⁴⁷. We’ll employ basic NLP (using libraries like spaCy or NLTK) to gauge the sentiment and relevance of news ⁵⁰ – for example, flag news that mentions our traded currencies or major market-moving terms (like “Fed surprise rate cut”). Each news item can be classified with an **impact score** (Low, Medium, High) based on keywords and possibly historical data (e.g. if similar news in the past caused a >50 pip move, classify as High) ⁵¹. When a **High-impact news** event is detected (whether scheduled or unscheduled), the NewsProcessor will issue a `NewsEvent` on the event bus, which the RiskEngine listens for. In response, the bot can **halt new entries immediately** (an *ImmediateKill* for unscheduled shocks) and even tighten stops or reduce positions on existing trades if appropriate ⁴⁹ ⁵². For example, if a sudden geopolitical conflict headline hits, the bot would stop opening trades at once ⁵³. We will log all such events and actions for audit (e.g. *“[News] Breaking: Fed emergency rate cut – trading paused”*) ⁵⁴. Integrating news awareness is crucial for FTMO-style trading because large unexpected moves can blow past stops. By avoiding trading during these times, we *reduce the risk of sudden rule violations due to slippage*. This integration essentially gives GENESIS a “fundamental awareness” complementing its technical strategy. We prioritize **ForexFactory for**

economic schedule (free and reliable calendar) and a **Bloomberg/Reuters feed for real-time news** (if budget allows – Bloomberg’s API is enterprise-grade). Even a **Twitter API or alternative** could be considered for certain breaking news (as some market-moving info hits social media first), though parsing Twitter adds complexity. The expected outcome is that GENESIS will intelligently **step out of the market during known high-risk events**, greatly reducing the chance of large slippage or unpredictable volatility hitting the account.

- **Alerting and Notifications (Telegram, TradingView, etc.):** To keep the human operator in the loop and to respond quickly to critical events, we will set up a robust alerting system. The bot will send real-time notifications for important states such as: kill-switch triggers (daily loss hit, etc.), errors/exceptions, execution problems (like order rejections or broker connection issues), and perhaps significant trade events (like a big winner or hitting a new equity high). We recommend integrating with **Telegram** for alerts, as it’s easy to push messages via a Telegram bot. For example, GENESIS can send a Telegram message “**ALERT:** Daily loss limit reached, trading halted for today.” the moment that happens. Similarly, on error it could send the exception summary. This gives the trader immediate insight if something is wrong, even if they aren’t watching the system UI. The Implementation Guide suggests using email/SMS or push notifications in such cases ⁵⁵ – Telegram is a modern equivalent with instant delivery. We can also configure **TradingView webhooks** as an input or output: for instance, if the user uses TradingView for technical analysis, custom TradingView alerts (via webhook URLs) could be sent to GENESIS to trigger specific actions or to notify the bot of a condition. Conversely, GENESIS could send certain alerts to a TradingView webhook that the user sets up, though usually it’s one-directional (TradingView -> bot) for signals. One concrete use case: the user draws a key support level on TradingView; when price hits it, TradingView can webhook that to GENESIS which might use it as a confluence signal. This kind of integration can be advanced, but it’s an option to **incorporate user-defined signals or levels** without modifying the bot code. At minimum, our plan is to implement a **Telegram Bot integration** and possibly email alerts as backup. All alert events will be published by the system (for example, the RiskEngine or NewsProcessor emits an “Alert” event with details), and a lightweight Alert service will catch those and forward to the configured channels. The user can also receive periodic summaries if desired (like end-of-day P/L report via Telegram). These alerting integrations ensure that **no critical event goes unnoticed** – even if the bot is running unattended, the trader will be pinged when something needs attention (e.g. the bot paused itself, or an account metric breached a threshold). It adds a layer of transparency and control expected in an institutional-grade system.

- **Trade Journaling and Analytics:** We will strengthen GENESIS’s trade journaling to facilitate both FTMO evaluations and internal strategy improvement. All trades and signals are already logged in detail; we propose to **centralize these logs and present them in a more user-friendly way**. For instance, we can integrate a database or use the existing event storage to record each trade’s data (entry, exit, P/L, duration, notes on why it was taken). A simple integration could be to push trades to a Google Sheet or a dedicated journaling tool, but since we have a web frontend, we will enhance the dashboard to include a **trade log view** and statistics. Each decision – from signal generation to risk check to execution – is being logged with a unique ID for traceability ⁵⁶. We will ensure these logs are stored long-term (e.g. in an Elasticsearch cluster or cloud storage) and can be queried for audit. Using the ELK stack is one recommendation ⁵⁷: all services log JSON-formatted events (with consistent fields like timestamp, service, event type, IDs) ⁵⁶, which are shipped to an Elasticsearch index. This would allow us to easily search and filter (e.g. “show all trades with confidence > 0.8 that hit SL”). It also aids performance monitoring (we can log execution latency, etc., and use Kibana or

Grafana to visualize). Additionally, we plan to track **trade performance metrics** for each strategy component – for example, win rate, average R:R ratio, drawdown, etc., which can be displayed on the dashboard or delivered in a daily report. Journaling is not only for compliance but for the trader's own improvement: GENESIS can effectively produce an **automatic trading journal** that the trader can review to see what patterns or conditions lead to wins or losses. We will also integrate a **daily equity curve and drawdown plot** on the dashboard, possibly overlaying FTMO limits, so one can visually monitor performance vs. the rules. In summary, by beefing up journaling we get both **accountability and insight** – every action is documented for FTMO or auditors, and the rich dataset can be mined to refine the strategy over time.

- **Multi-Account Orchestration:** A key upgrade is enabling GENESIS to handle **multiple trading accounts simultaneously** – for example, managing several FTMO challenge accounts or funded accounts in parallel, or splitting capital across broker accounts. The architecture's event-driven nature makes this feasible: we can either run multiple instances of the bot (one per account) or run one unified system that publishes signals to a topic and has multiple Execution/Risk consumers (one per account) subscribing to those signals. We propose the latter for efficiency. We will introduce an account abstraction layer where each ExecutionEngine instance is tied to a specific account (with its MT5 login credentials) but they all listen to the same `trade_signals` topic. For example, if a BUY signal on EURUSD is generated, **all active account engines receive it and decide to act**. They could each trade it, possibly with different sizing or risk settings per account. This effectively provides a built-in trade copier – a signal from a master strategy can be executed across multiple accounts at once. Because our services are decoupled, adding a new account is as easy as launching another ExecutionEngine (and RiskEngine) service in a new consumer group on the bus ⁵⁸. The new service will subscribe to the existing event topics (market data, signals, etc.) and operate independently on its account ⁵⁹. We'll need to ensure that account-specific risk limits are tracked separately by each instance (since each account has its own 5%/10% limits, etc.) – that's straightforward by parameterizing the RiskEngine with the account's balance and rules. We should also segregate any account-specific state like open trades; that stays internal to each Execution/Risk instance. If desired, we can also stagger strategies: e.g. run slightly different parameters on different accounts (account A might take only subset of signals or smaller size) by customizing the RiskEngine or filter settings per instance. Communication-wise, all account instances can still report to one central dashboard, tagging their logs with an account ID. The dashboard could then display either aggregate or per-account views. For example, multi-account equity curves, or trade logs filtered by account. From a deployment perspective, we can use Docker Compose or Kubernetes to spin up N copies of the necessary services for N accounts, with configuration differences handled via environment variables (for login, etc.). We might also implement a **master account orchestrator** service that can broadcast commands – e.g. "halt trading on all accounts now" or rotate credentials if needed. But fundamentally, the event-driven design makes multi-account support quite elegant: it's a publish-subscribe model where adding a subscriber (account) doesn't require changing the publisher (strategy). The expected outcome is that GENESIS will **scale to manage multiple FTMO accounts** (or other broker accounts) concurrently, enabling higher combined capital to be traded with the same strategy. This is valuable for a trader managing prop firm allocations or trading a personal and a funded account in tandem. It also adds redundancy – if one account has an issue (e.g. reaches a limit or gets suspended), the others can continue. Overall, multi-account orchestration will transform GENESIS from a single-account bot into a **mini portfolio management system** capable of trade replication and coordinated risk management across accounts.

4. Containerized Deployment Architecture (Single & Multi-Account)

To support the above improvements and ensure easy deployment, we recommend a fully **containerized architecture** using Docker (and optionally Kubernetes for orchestration). Each microservice – Market Data Feed, Signal Engine, Pattern Recognizer, NewsProcessor, ExecutionEngine, RiskEngine, Dashboard – will run in its own container, communicating over the event bus and APIs. This aligns with the current design where services can be deployed independently ⁴. Containerization offers consistency across environments (dev, test, prod) and isolates dependencies for each module. We can define a Docker Compose configuration that brings up the entire suite of services, including infrastructure like Kafka/Redis, and links them together.

Event Streaming Backbone: We will deploy **Kafka** (or a Redis Stream) container as the central message broker. Kafka is preferred for its durability and scaling. All services will connect to Kafka to publish/subscribe topics (e.g. `ticks`, `trade_signals`, `news`, `orders`, `alerts`). We will tune Kafka for low-latency messaging (e.g. using smaller batches and enabling immediate flush for critical topics) to ensure real-time delivery. If using Redis Streams as a simpler alternative, it could also work for moderate volumes, with perhaps less built-in persistence than Kafka. Either way, this decoupled event bus is the spine of the system, enabling **asynchronous, loosely-coupled communication** ¹⁰. One benefit is easier **horizontal scaling** – if one service becomes a bottleneck, we can scale it out by adding more consumers to its group ³. The containers can even be distributed across multiple servers if needed, as long as they connect to the central bus.

Containerization for Multi-Account: As noted, for multiple accounts we might run multiple instances of certain services. In Docker, this can be handled by service replication. For example, we could label ExecutionEngine+RiskEngine for account1, account2, etc., each as separate service in the compose file (or scale a service and pass account-specific env vars). Kubernetes would make this even more dynamic with replicas and configmaps per account. Each instance would join the appropriate consumer group so they don't steal events from each other (or use key-based partitioning by account). In Kafka, we can design the topics such that signals are broadcast to all (using consumer groups where each account service group gets a copy of signals) ⁵⁸. Alternatively, we could partition topics by account if we wanted to segregate strategies, but likely they share the same signal stream.

Orchestration & DevOps: We suggest using **Kubernetes** in the long run for its self-healing and scaling capabilities. Kubernetes can monitor the health of each container (via liveness probes) and restart any failed pods automatically. This complements our in-app watchdogs. We can also use Kubernetes to do rolling updates service by service, which is useful for quick strategy tweaks with minimal downtime. In a simpler setup, Docker Compose on a single server can manage the containers with restart policies (`always` restart to handle crashes). A cloud deployment (AWS ECS/EKS, Azure AKS, etc.) could be considered if needed for reliability. The containers themselves will be designed to be stateless (except for any local caches) – all state either goes to the database/event-log or in-memory. This makes scaling and restarting safe.

State Management: We will include a small **database container** (e.g. PostgreSQL or MongoDB) if needed for persistent state and configuration (for example, storing strategy parameters, recording completed trades for long-term, or saving pattern recognition models). However, many aspects might be covered by the event log (Kafka) itself as an "event store". Still, an SQL/NoSQL DB can be handy for quick queries or the dashboard's needs. We'll ensure any such DB is also run redundantly or with backups.

Monitoring & Logging Infrastructure: We will containerize our logging/monitoring stack as well. For instance, include **Elasticsearch and Kibana** containers (or use a cloud logging service) to collect logs from all services. A **Prometheus** container can scrape metrics from the services (each service can expose an HTTP `/metrics` endpoint or push metrics) ⁶⁰ ⁶¹, and a **Grafana** container can display dashboards. This gives a professional ops setup where we can see at a glance the system health: e.g. feed latency, queue lengths, CPU/memory of each container, number of signals/hour, P/L curve, etc. ⁶⁰ ⁶². We'll also set up alerting rules in Prometheus (like if no ticks for 60s, or if latency spikes, etc., raise an alert event) ⁶².

Security & Configuration: Each container will have its own config, with sensitive info (API keys, account passwords) passed via environment variables or mounted secret files (never hard-coded). Inter-service communication will occur either inside a Docker network (which is isolated) or over secured channels (Kafka can be set up with SSL, etc.). We will ensure that even if the containers are compromised, minimal information is exposed in logs (e.g. we won't log plaintext passwords) ⁶³. For FTMO and compliance, we might also consider encrypting log storage or at least protecting it, since trade data could be sensitive ⁶³.

In summary, the containerized architecture will mirror the conceptual microservices design with an event bus in the middle. An **example deployment** might consist of:

- **Broker Bridge (MT5 API)** – Python service connecting to MetaTrader 5 terminal (this runs on a Windows server or via Wine on Linux, but the Python MT5 API can run in a container if it has network access to the MT5 terminal). Alternatively, we might run the MT5 terminal itself on the host and have the Python connect via TCP. This part needs careful setup due to MT5's dependency on a client terminal. Possibly we use a Windows VM for MT5 and the rest on Linux containers.
- **MarketDataFeedManager** – Python container that logs in to MT5 and subscribes to ticks for all instruments, publishes to Kafka topic `ticks` ⁶⁴ ⁶⁵.
- **SignalEngine** – Python container that subscribes to `ticks` (and any other event topics like patterns or news) and produces trade signal events to `trade_signals`. This includes technical indicator logic and multi-timeframe analysis.
- **PatternRecognition service** – (if separate) Python container that subscribes to `ticks` (or bar data) and identifies chart patterns, publishing `pattern` events with confidences.
- **NewsProcessor** – Python container that calls news APIs and publishes `news` or `kill_switch` events when needed ⁴⁷.
- **RiskEngine** – Python container that subscribes to `trade_signals` and `news` events, evaluates risk constraints. It can either act as a gatekeeper (only forward allowed signals to ExecutionEngine) or just flag issues. Our design likely has Risk checks inside ExecutionEngine process for simplicity, but we can have a dedicated risk microservice that the ExecutionEngine queries.
- **ExecutionEngine** – Python container (one per account) that subscribes to `trade_signals` (filtered or via risk) and places orders on MT5 via the API, also publishing `order_status` events (fills, etc.). It also listens for a global `pause_trading` flag or kill-switch events to know when to suspend trading ⁶⁶.
- **Dashboard** – Node.js/React container serving the web UI. It may subscribe to certain Kafka topics (through a websocket bridge or by an API that aggregates data) to show live updates (e.g. via Socket.io as in the Heroku reference ⁶⁷ ⁶⁸).
- **Alert service** – Small Python/Node container that subscribes to `alerts` topic and forwards to Telegram/email.
- **Logging/Monitoring** – Elasticsearch, Kibana, Prometheus, Grafana containers as needed, plus perhaps Filebeat or Logstash to pipe logs.

This architecture not only supports **single vs. multiple accounts** by scaling Execution/Risk containers, but also is **cloud-ready and fault-tolerant**. If one container dies, it's isolated (others continue running) ³, and the orchestrator will revive it. If needed for high availability, we can run multiple instances of critical services and even multiple brokers (though FTMO typically one account at a time per instance). Kafka can be clustered across nodes to avoid any broker downtime ⁹. All of these steps make the deployment *institutional-grade*: easy to update (you can deploy a new version of, say, the SignalEngine container without touching others), easy to maintain, and scalable on demand.

In effect, we are adopting a cloud-native deployment for GENESIS, suitable for professional trading operations or prop firm usage. The containerized multi-account setup allows an operator to manage a "fleet" of accounts through one coherent system, which is a strong advantage in the FTMO/funded account arena.

5. Code Quality, Modularity, and Audit-Readiness (10x Standard)

We will conduct a thorough audit of the codebase and implement improvements to reach a "10x" engineering standard in terms of clarity, reliability, and maintainability. Key focus areas:

- **Strict Modular Design (SOLID Principles):** We will ensure the code reflects the separation of concerns in the architecture. Each module/service in code should have a single responsibility and a well-defined interface ⁶⁹. For example, the Signal generation code shouldn't directly place trades or access account info – it should only output signals. The Execution code shouldn't compute technical indicators – it should only handle orders. This modularity makes the code easier to test and reason about. We'll refactor any overly entangled parts. For instance, if the current implementation has one big script handling everything, we will break it into classes or components (FeedManager, SignalEngine, etc.) as described in the Implementation Guide. This also extends to configuration – using config files or objects to pass parameters rather than hardcoding values deep in functions. Following these practices will produce a **clean, well-organized codebase** where adding or modifying features can be done in isolation.
- **Robust Error Handling & Logging:** As mentioned, we'll wrap critical sections with try/except and make liberal use of logging. Every significant event (tick received, signal generated, trade executed, risk check failed, etc.) should be logged with context ⁴⁶. We will use Python's `logging` library configured to output structured JSON logs (with fields like timestamp, service, event type, details) ⁵⁶. This structured logging means the data can be parsed and analyzed easily (for example, counting how many signals had confidence > 0.8, etc.) ⁵⁶ ⁷⁰. We'll include correlation IDs to tie together events across services – e.g. a signal ID that gets carried through to the order and outcome – which makes debugging and auditing much easier ⁷¹. On error handling: we'll ensure that any exception is caught at least at the top level of each service loop, and logged with stack trace. We won't let the service just crash silently. If an error is non-critical, the service can log it and continue. If it's critical (e.g. cannot connect to broker), the service can notify via an alert and perhaps retry or shut down gracefully. Importantly, if something truly unexpected happens (bug), the system should **fail safe** – e.g. pause trading rather than continue in a possibly inconsistent state ¹⁷ ¹⁸. The code will include checks for common failure modes: e.g. if we don't get a confirmation for an order placed, if data feed stops, etc., and handle them. All these logs and error catches contribute to an **audit trail** – we want to be able to trace every decision and outcome after the fact, and also have the system surface issues in real-time.

- **Performance Profiling and Optimization:** We will profile the code to identify any slow spots or inefficiencies. Given Python's GIL, one area to check is that no single thread is doing heavy computations that block others. For example, if the SignalEngine does heavy indicator math on a single thread, it might lag when many ticks arrive. We can use asynchronous programming or background threads to offload tasks (e.g. compute indicators on historical data in one thread while main thread processes new tick events). If necessary, we might leverage Python's `multiprocessing` or an in-memory job queue to distribute workload. We'll also ensure that we use efficient data structures (e.g. NumPy arrays for numeric calc rather than Python lists in loops). Another angle is to implement critical path functions in Cython or use vectorized libraries (pandas, talib, etc.) for speed. The expected outcome is that after optimization, the system can handle the load (perhaps defined as X ticks per second across Y symbols) without falling behind. We'll include **metrics collection** in code, such as measuring how long each signal processing step takes, or how long between a signal and order placement. These can be logged or exposed to Prometheus ⁶². If we notice any unacceptable delays, we revisit the code. Memory usage will also be monitored to ensure no leaks (long-running bots need stable memory). Essentially, our code audit will treat performance as a first-class concern along with correctness.

- **Signal Processing and Confluence Architecture:** We will refine the signal-generation logic to use a **multi-factor confluence** approach for higher accuracy. This means that a trade signal should ideally be backed by multiple independent indicators or confirmations (trend, momentum, pattern, etc.) ⁷² ⁷³. The code should be organized to combine these inputs cleanly – for example, computing various indicator values and then checking a set of conditions for entry. If not already, we'll introduce a **confidence score** for each signal that quantifies its strength. The Implementation Guide provides an approach to compute a confidence value based on how many conditions line up and perhaps weighting them ⁷⁴. We will implement such scoring, e.g. on a 0 to 1 scale. Moreover, we plan to incorporate **adaptive confidence adjustment**: the bot can learn over time which signals are more reliable. For instance, track the performance of each pattern or indicator signal historically, and increase/decrease its weight in the confidence calculation accordingly ⁷⁵ ⁷⁶. If a certain setup has yielded 70% win rate historically, give it high confidence; if another has 40%, treat it skeptically. Over time this becomes a feedback loop (part of an adaptive learning engine). Code-wise, this means maintaining statistics (counters of wins/losses for each signal type) and updating parameters periodically. We'll also ensure that the **ensemble logic** for combining signals is sound – possibly using a weighted average or rules that require multiple green lights before a trade is taken ⁷³. The outcome should be **fewer but higher-quality trades**: the bot waits for "A-grade" setups where there is confluence (e.g. trend is up, momentum indicator bullish, a bullish chart pattern present, and no conflicting news). This increases the probability of success per trade, which is crucial in FTMO where you have limited drawdown. We will unit test the signal logic extensively – feed historical data and verify it triggers signals only when expected. Also, log each signal with the factors that were true/false for transparency ⁷⁷.

- **Adaptive Strategies and Continuous Improvement:** As part of 10x standards, we don't want the strategy to be static. The codebase should be structured to allow easy strategy tweaks and tuning. We plan to implement (if not already) an **Adaptive Learning module** that can adjust strategy parameters based on recent performance (likely offline or on weekends as a batch process) ⁷⁸ ⁷⁹. This could involve walk-forward optimization or even reinforcement learning to find better settings (e.g. optimal stop-loss distance, or which hours of day are best to trade) ⁷⁹ ⁸⁰. While this is an advanced feature, laying the groundwork in code (like reading key parameters from a config file and

allowing them to be updated) is important. That way, improvements can be deployed without code changes (just config updates) and possibly automated. We will also maintain **version control and testing** rigorously – every code change will be tested against historical data to ensure it would not have broken rules or caused big losses ⁸¹. We want to see that our improvements truly improve performance (for example, test that after adding news filter, the bot avoids known news-related drawdowns in history). By logging everything and testing, we uphold an audit standard where any strategy change is documented and can be justified with data (this is good practice especially if presenting to investors or firm).

- **Compliance and Audit Trails:** In a funded account scenario, you may be asked to prove your trades were according to plan. Our code will maintain an **immutable audit trail** of decisions. We will store logs securely (write-once or with tamper-evident measures like hashing log files) ⁸². If needed, we might implement an audit export that compiles all trades and their justifications (from the logs) for a given period. The goal is that an external auditor could verify that, for example, no trade was taken without SL, no day had >5% loss, etc., simply by reading our records. We'll also reconcile the bot's records with broker statements regularly to catch any discrepancies ⁸³. All these procedures elevate the system to **institutional compliance level**.

In summary, the code quality upgrades will result in a **clean, maintainable, and well-documented codebase**. Each component will be clearly defined and tested. The system will produce rich logs that not only help in debugging but also form an audit record. By enforcing best practices and thorough testing, we reduce the chance of hidden bugs that could cause costly errors in live trading. The end result is a bot that a professional firm would be comfortable running, knowing it's been built and audited to a high standard, with every edge case considered and every rule enforced.

6. Current Limitations & Recommended Enhancements

In our audit, we identified some current limitations of the GENESIS bot implementation. Below we summarize these issues and our recommendations (many of which have been detailed above) to enhance the system using industry best practices:

- **Limitation 1: Potential Single-Threaded Bottlenecks in Feed Handling.** The current Market Data feed may be using a single-threaded loop to process ticks for all instruments, which could struggle during peak market activity. **Enhancement:** Implement asynchronous or multi-threaded feed handling. By using non-blocking I/O or dedicated threads for reading ticks, and possibly partitioning symbols across threads, the feed can ingest data with minimal latency ⁷. This ensures the bot keeps up with fast markets (e.g. high-frequency tick bursts) without lag, maintaining data fidelity for the strategy.
- **Limitation 2: Use of Market Orders (Slippage Risk).** If the current version ever uses market orders for entry/exit, it risks uncontrolled slippage which could violate FTMO limits or cause inconsistent execution. **Enhancement:** Strictly switch to limit orders for all entries and exits (no market orders) ²⁰. As discussed, this sniper execution style guarantees price or no fill, dramatically reducing slippage. Any existing code paths that call `market_order()` will be replaced with limit logic, including smart re-quote and partial fill handling. The benefit is tighter execution control and improved compliance with risk limits.

- **Limitation 3: Incomplete FTMO Rule Enforcement.** The earlier implementation might not have fully enforced the 5% daily loss or 1:30 leverage rules in real-time. For example, it might rely on the trader to stop it rather than an automatic kill-switch. **Enhancement:** Embed all FTMO rules as hard constraints in code. We introduced automated daily loss tracking with preemptive halt at ~4.5% and full stop at 5% ⁸⁴ ³¹. We also enforce total drawdown halts at 10% ³⁵, leverage checks before every order ³⁷, and mandatory stop-loss on every trade ⁴¹. These measures ensure the bot **cannot** break the rules – an order that would violate a rule is simply not executed. Implementing these was top priority to protect the account and pass FTMO evaluations.
- **Limitation 4: Lack of News/Event Risk Management.** The current bot likely does not integrate any news feeds, so it keeps trading through high-impact news or unpredictable events, which could cause large losses or slippage. **Enhancement:** Integrate economic calendar and news feed processing. By pausing trading during known major news (e.g. FOMC, NFP releases) and halting on surprise news ⁴⁷ ⁴⁹, the bot avoids the most dangerous market conditions. This reduces the likelihood of hitting loss limits due to a single spike. We also widen stops or avoid new trades around these times to manage volatility. This enhancement brings human common-sense (any experienced trader is cautious around news) into the automated system.
- **Limitation 5: Limited Order Execution Adaptiveness.** The prior implementation may execute orders in a fire-and-forget manner (e.g. send a market or limit order and not react if not filled or partially filled). That can lead to missed trades or suboptimal fills. **Enhancement:** Introduce adaptive order logic – monitor orders after placement and react. For instance, if a limit order isn't filled within X seconds, consider moving the price (chasing a little) or canceling if the move has run away ⁸⁵. If an order is partially filled, implement logic to fill the remainder via new orders or decide to skip the rest if conditions changed ²⁴. Essentially, the ExecutionEngine becomes smarter, emulating how a human trader might adjust an order. The benefit is a higher likelihood of getting into good trades (with controlled adjustments) and not getting stuck with half-fills or missed opportunities.
- **Limitation 6: Confidence and Position Sizing Not Dynamic.** The current strategy might be using a fixed threshold for entries and uniform position sizing regardless of signal quality. This can be improved. **Enhancement:** Incorporate a **confidence scoring system** for signals and adjust position size or decision threshold based on it. Using the methods described (multi-factor confluence, pattern reliability stats, etc.), the bot can assign a confidence 0-1 to each signal ⁷⁴ ⁸⁶. We then set a rule that only signals above a certain confidence (say 0.7) trigger trades, or we vary the risk per trade proportional to confidence (within limits). Moreover, maintain performance stats to recalibrate these confidences – if some signals consistently fail, lower their weight; if some are highly accurate, increase their impact ⁷⁵ ⁷⁶. This adaptive confidence approach means over time the bot **learns which setups are truly high probability**, leading to better risk-adjusted performance. We expect this to increase win rate and/or average R:R, which helps meet FTMO profit targets more safely.
- **Limitation 7: Monolithic or Tightly Coupled Codebase.** If the provided code (Genesis FINAL TRY) is not fully modular – for example, if it mixes trading logic with UI or doesn't have clear separation of functions – that makes it harder to maintain and extend. **Enhancement:** Refactor the code into a modular structure reflecting the microservices design. Each component (data feed, signals, execution, risk, etc.) should be encapsulated, even if running in one process. Use classes or at least distinct modules for each. Ensure communication is via clear interfaces/events rather than direct function calls that create hidden dependencies. This also allows easier unit testing of each part.

Adopting this structure will make future upgrades (like adding a new indicator or switching brokers) much simpler. It also aligns with our containerization plan – code structure should map to deployable units.

- **Limitation 8: Insufficient Logging and Transparency.** The current implementation might not log enough detail or in a structured way, which could hinder debugging and audit. For instance, it might print some info to console but not store historical logs, or not log the rationale behind trades. **Enhancement:** Implement comprehensive structured logging (as described in Section 5). Every tick, signal, decision, order, and result should produce a log entry with machine-parsable data ⁵⁶ ⁴⁶ . Store these logs centrally so we can analyze them later. Additionally, improve the **dashboard journaling**: show a table of all trades with entries, exits, P/L, and maybe notes on why they were taken. Possibly include a feature where clicking a trade shows the indicators state at that time (since we log those). This level of transparency is key for a “professional-grade audit.” It means any anomalies can be traced to root cause quickly, and the trader (or FTMO) can verify the bot is behaving as intended.
- **Limitation 9: Deployment and Environment Setup.** Currently, running the bot may require manual steps (starting MT5 terminal, running the script, etc.), which is prone to error. **Enhancement:** Move to the containerized deployment using Docker, so the environment is reproducible. Use docker-compose to define all services, so with one command the entire system comes up configured. This reduces setup time and errors like version mismatches. It also paves the way for deploying in the cloud or on remote servers easily. The outcome is a **consistent runtime environment** wherever the bot is deployed, and easier collaboration (others can run the same setup from the repo). Moreover, containerization combined with Git version control means we can roll back to previous stable versions quickly if an issue is found – useful for a mission-critical trading system.
- **Limitation 10: Testing and Validation.** It’s unclear if the current iteration has thorough automated tests or if it was mainly tested manually. Lack of testing can allow bugs to slip through or rules not properly implemented under all scenarios. **Enhancement:** Develop a robust testing harness for GENESIS. This includes unit tests for each module (e.g. test that the RiskEngine rejects trades correctly when daily loss would exceed 5%, etc.), and integration tests using historical data replay. We can simulate one day of tick data through the system and verify that the outcomes (trades taken, P/L) match expected results. We will specifically test edge cases: e.g. a scenario with a huge gap at market open to see if the bot closes trades that exceed limits (even if slippage caused an over-loss, how does it respond?) ⁸⁷ . Another test: feed consecutive losing trades quickly to ensure the daily kill-switch triggers at the right point ⁸⁸ . Testing gives us confidence (and is part of the 10x standard) that the bot will behave correctly in live markets, and it also guards against future regressions when we modify code. Expected outcome is fewer runtime surprises and a system that does exactly what’s written in the spec.

By addressing these limitations with the recommended enhancements, we align GENESIS with **industry best practices** for algorithmic trading. Each improvement – from async feeds to dynamic risk adjustments – contributes to a more robust and profitable trading bot. The cumulative effect is a system that is *safer, faster, and smarter* than before, well-suited for the stringent FTMO environment and for scaling up to larger capital.

Prioritized Action Plan and Expected Outcomes

Based on the audit findings and proposals above, we outline a step-by-step action plan to implement these upgrades. The steps are ordered by priority (considering risk impact and foundation laying), along with the expected outcomes for each:

- 1. Implement FTMO Risk Kill-Switches and Order Constraints (Immediate Priority):** *Action:* Modify the RiskEngine/ExecutionEngine to enforce the 5% daily loss halt, 10% max drawdown stop, 1:30 leverage check, and “no SL, no trade” rule. Also switch all entries to limit orders only. *Outcome:* Instantly eliminates the possibility of catastrophic losses or rule violations. Even if other parts of the bot have issues, these safeguards protect the account from blowing up ³⁰ ³⁷. The bot will abide by FTMO rules 100% – improving trust and guaranteeing challenge eligibility is not lost by a technicality.
- 2. Enhance Logging & Alerting (Immediate Priority):** *Action:* Introduce detailed structured logging across all modules and set up the alerting system (Telegram/email notifications for critical events). *Outcome:* Within days, we'll have full visibility into the bot's actions. Any error or rule trigger will ping the human operator immediately ⁸⁹, allowing for quick intervention if needed. The detailed logs mean that if something goes wrong, we can diagnose it from the log trail. This also prepares us for later steps like analysis and confidence tuning, since we'll be collecting rich data from the start.
- 3. Refactor Code into Modular Microservices:** *Action:* Restructure the codebase to isolate components (feed, signals, risk, execution, etc.) and define clear event interfaces between them. Remove any redundant or tightly-coupled code. *Outcome:* A cleaner, more maintainable codebase ready for extension. This refactoring likely won't change functionality immediately but sets the stage for all other improvements and containerization. It reduces development friction and lowers the chance of bugs (since each part can be understood and tested in isolation) ⁶⁹.
- 4. Deploy Kafka/Redis Event Bus and Containerize Services (High Priority):** *Action:* Introduce Kafka or Redis Streams as the messaging layer in the development environment. Adjust services to use it (if not already). Create Dockerfiles and a docker-compose setup for all components including Kafka. *Outcome:* The system will now run in a reproducible containerized form, and communications will be truly decoupled via the event bus. This will slightly increase complexity, but it yields **scalability and reliability** – e.g. if the Signal service crashes, Kafka buffers events until it restarts ⁹. Containerization also means we can deploy to a server or cloud easily and run multiple instances as needed.
- 5. Integrate NewsProcessor and Economic Calendar (High Priority):** *Action:* Develop the NewsProcessor service to fetch economic calendar events (ForexFactory) and real-time news (e.g. a free news API or RSS feeds). Implement logic to parse and classify news, and generate pause/resume signals to the trading system ⁴⁷ ⁴⁹. Hook this into the RiskEngine (or directly into ExecutionEngine) to actually pause trading during events. *Outcome:* Within a short time, the bot will start **avoiding trades around major news releases** and will react to breaking news by protecting open trades or stopping new ones. We expect to see a reduction in volatile whipsaw trades and more stable performance during normally perilous times. This step materially decreases the likelihood of a sudden large drawdown (which could violate FTMO daily loss) due to an unforeseen headline.

6. **Sniper Execution Logic & Partial Fill Handling (High Priority):** *Action:* Overhaul the ExecutionEngine's order handling as described – placing limit orders with a few pips offset ²¹, monitoring for fills, adjusting or canceling orders that don't fill in time, and implementing partial fill retries ⁸⁵ ²⁴. Also incorporate time-of-day filters (no trading during illiquid hours). *Outcome:* The execution quality of trades will improve noticeably. We anticipate better entry prices on average (since we often get price improvement via limit entries) and fewer instances of slippage. Partial fills will be managed so that position sizes are as intended or gracefully reduced – no more hanging orders. Overall, execution will feel more “professional” and less like a basic EA; it will seek optimal fills like a human trader would. This should boost net performance and consistency.
7. **Multi-Account Support (Medium Priority):** *Action:* Extend the system to handle multiple accounts. This involves allowing multiple ExecutionEngine instances and segregating risk management per account. Test by connecting to two demo accounts in parallel and copying signals to both. *Outcome:* GENESIS can now multiply its impact by trading several accounts at once with the same signals. For the user, this means the ability to scale up (e.g. if you have 2-3 funded accounts, the bot can manage all simultaneously) or to diversify (maybe one account trades a subset of instruments, another a different subset). The immediate result is more flexibility; strategically, it allows growth of trading operations without changing the core strategy. It also adds redundancy – one account could fail or be in drawdown while others might still be fine, smoothing overall equity.
8. **Confidence-Weighted Position Sizing and Filtering (Medium Priority):** *Action:* Introduce the confidence scoring system for signals. Use historical backtests or live forward-testing data to calibrate an initial confidence for each strategy condition, and implement logic to only trade high-confidence signals or to size them proportional to confidence. *Outcome:* Fewer low-quality trades will be taken – expect a dip in trade frequency but an increase in win rate or average profit per trade. Over time, as the bot learns, the confidence should correlate with actual outcomes, making the bot increasingly selective in a good way ⁷⁵ ⁹⁰. This improves the **risk-adjusted return** of the strategy, helping to meet profit targets with less volatility – a key to passing FTMO challenges and maintaining funded accounts. It also provides a transparent metric (confidence) that can be monitored; for instance, the dashboard can show “Current signal confidence: 0.85” which gives the trader extra insight.
9. **Automated Testing Framework (Medium Priority):** *Action:* Develop a suite of tests for critical functions. Implement a backtest mode for the bot where it can consume historical tick data from a file or database (perhaps by writing a replay script that reads stored ticks and publishes them to Kafka at high speed). Simulate various scenarios: big gap opens, consecutive losses, news events, etc. Ensure the bot's responses match expectations (like triggering kill-switches appropriately). *Outcome:* A robust validation of the system before each release. This will catch any regression where, say, a refactor accidentally disables a risk check or a logic bug that could surface only in unusual conditions. Having tests also speeds up development – we can change code and quickly run tests to ensure nothing broke. This **increases reliability in live trading**, as we've essentially run many “what-if” scenarios in advance. FTMO challenges are time-limited (e.g. 30 days); an unnoticed bug could ruin a challenge in one day, so testing is critical to avoid that. With this step, we gain high confidence that the bot will behave correctly in the wild.
10. **Deployment on a Stable Hosting Environment (Low/Continuous Priority):** *Action:* Once all the above are implemented and tested on a local setup, deploy the containerized system to a reliable

server (VPS or cloud instance). Set up monitoring (Grafana dashboards, alerting on server resource issues). Possibly employ a process manager or Kubernetes on the server for auto-restart and uptime. *Outcome:* GENESIS will run 24/7 with minimal downtime. If anything goes wrong, the team is alerted immediately via the earlier alerting mechanism. The deployment is reproducible, so moving to a new server or scaling out is straightforward. Essentially, the bot transitions from a “development project” to a **production-grade trading system**.

Each of these steps builds on the previous, and many can be done in parallel by a team. By following this prioritized plan, we ensure that the most critical risk protections are in place first, then we enhance the bot’s intelligence and performance, and finally polish the deployment and testing aspects. The expected end result after completing these steps is that **GENESIS will be transformed into an institutional-quality trading platform**: one that is fast, resilient to errors, compliant with strict risk rules, adaptive to market conditions, and thoroughly monitored and audited. This sets the stage not only for excelling in FTMO challenges but also for long-term profitable trading in funded accounts or even managing private capital at scale. Each improvement directly contributes to either *protecting capital* or *improving returns* – ultimately increasing the probability of success under FTMO’s evaluation and beyond.

1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30
31 32 33 34 35 36 37 38 39 40 41 42 43 44 45 46 47 48 49 50 51 52 53 54 55 56 57 60 61 62
63 64 65 66 69 70 71 72 73 74 75 76 77 78 79 80 81 82 83 84 85 86 87 88 89 90 GENESIS

Trading Bot Implementation Guide.pdf

file:///file-9ciUcMZzLSbkkUionY6fcL

58 59 67 68 Reference Architecture: Event-Driven Microservices with Apache Kafka | Heroku Dev Center

<https://devcenter.heroku.com/articles/event-driven-microservices-with-apache-kafka>